

async-fl

February 13, 2023

0.1 Async Training

The goal here is to compute the average value of all the client. Each client has a different compute speed than the other clients, depending on the ‘compute time’ distribution (Normal, Exp, Uniform). This toy example simulates an asynchronous (federated) learning system. With some extensions, this should be able to run as a simulation on a GPU.

Author: Bart Cox 13-02-2023

0.2 Import dependencies

```
[ ]: import numpy as np
      import matplotlib.pyplot as plt
      import seaborn as sns
      import pandas as pd
      from tqdm import tqdm
```

0.3 Aux functions

```
[ ]: def create_client(id: int, ct=1):
      return {
          'value': np.random.randint(0, 1000),
          'id': id,
          'picked': 0,
          'age': 0,
          'ct': ct,
          'ct_left': ct,
      }

def update(central_model, client):
    staleness = central_model['age'] - client['age']
    central_model['age'] += 1
    central_model['value'] += client['value']

    central_model['approx'] = central_model['value'] / central_model['age']
    central_model['difference'] = abs(central_model['true_avg'] - ↴
        central_model['approx'])
```

```

client['picked'] += 1

client['age'] = central_model['age']
return central_model, client, staleness

```

[]: def run_with_compute_times(num_clients, num_rounds, compute_times):
 clients = [create_client(x[0], x[1]) for x in zip(range(num_clients), ↴compute_times)]
 # Calculate the target value
 avg_value = np.mean([x['value'] for x in clients])
 central_model = {
 'value': 0.0,
 'age': 0,
 'approx': 0,
 'difference': 0,
 'error_abs': 0,
 'true_avg': avg_value
 }

 # print('Training')
 data = []
 for _round in tqdm(range(num_rounds)):
 rc = min(clients, key=lambda x:x['ct_left'])
 # Find client the finishes first
 min_ct = rc['ct_left']

 # Update the time for all the clients with the elapsed time of min_ct
 for c in clients:
 c['ct_left'] -= min_ct

 # Perform client server interaction between server and min_ct
 central_model, rc, staleness = update(central_model, rc)
 # Reset compute time of min_ct
 rc['ct_left'] = rc['ct']
 clients[rc['id']] = rc

 # Gather some statistics
 data.append([central_model['age'], central_model['approx'], 'approx', ↴staleness])
 data.append([central_model['age'], avg_value, 'avg', 0])
 data.append([central_model['age'], abs(avg_value - ↴central_model['approx']), 'error_abs', 0])
 return pd.DataFrame(data, columns=['age', 'value', 'type', 'staleness'])

def run(num_clients = 100, num_rounds = 1000):

```

results = []
# Normal distribution
mu, sigma = 15, 3
ct_normal = np.random.normal(mu, sigma, num_clients)
df_normal = run_with_compute_times(num_clients, num_rounds, ct_normal)
results.append([
    f'Normal distribution (mu={mu}, sigma={sigma}), client={num_clients},',
    f'rounds={num_rounds}', ,
    ct_normal,
    df_normal
])

# Exponential
exp_val = 1
ct_exp = np.random.exponential(exp_val, num_clients)
df_exp = run_with_compute_times(num_clients, num_rounds, ct_exp)
results.append([
    f'Exponential distribution (lambda={exp_val}), client={num_clients},',
    f'rounds={num_rounds}', ,
    ct_exp,
    df_exp
])

# Uniform
uniform_vals = [1, 20]
ct_uniform = np.random.uniform(*uniform_vals, num_clients)
df_uniform = run_with_compute_times(num_clients, num_rounds, ct_uniform)
results.append([
    f'Uniform distribution (min-max={uniform_vals}), client={num_clients},',
    f'rounds={num_rounds}', ,
    ct_uniform,
    df_uniform
])

# plot results
for name, ct, df in results:
    fig, ax = plt.subplots(nrows=1, ncols=4, figsize=(20, 5))
    fig.suptitle(name)
    ax[0].set_title('Generated Compute times')
    df_ct = pd.DataFrame(ct, columns=['ct'])
    sns.histplot(ax=ax[0], data=df_ct, x='ct')
    ax[0].set_xlabel('Compute time')
    # sns.histplot(ax=ax[0], data=df_ct, x='ct', bins=20)
    ax[1].set_title('Value approximation')
    local_df = df[df['type']!='error_abs']
    sns.lineplot(ax=ax[1], data=local_df, x='age', y='value', hue='type')
    ax[1].set_xlabel('Model age')

```

```

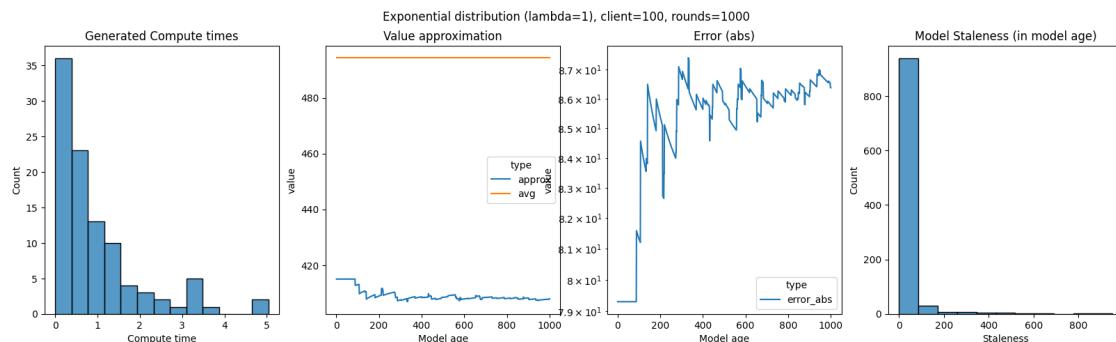
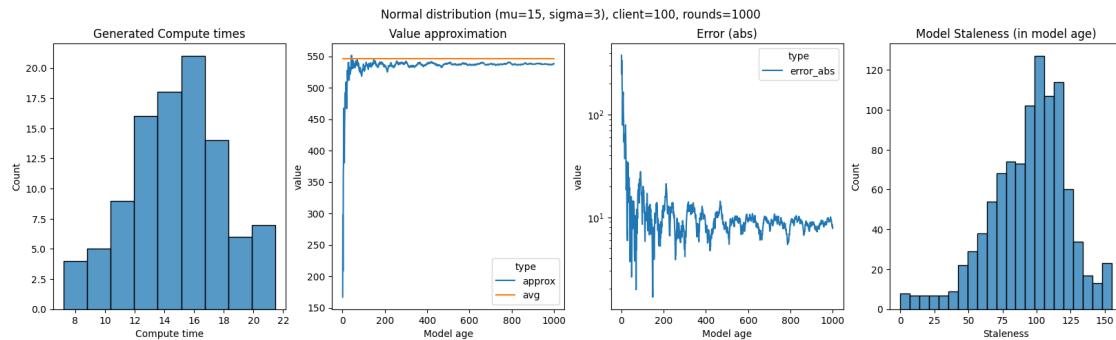
ax[2].set_title('Error (abs)')
local_df = df[df['type']=='error_abs']
sns.lineplot(ax=ax[2], data=local_df, x='age', y='value', hue='type', )
ax[2].set_xlabel('Model age')
ax[2].set_yscale('log')
ax[3].set_title('Model Staleness (in model age)')
local_df = df[df['type']=='approx']
sns.histplot(ax=ax[3], data=local_df, x="staleness")
ax[3].set_xlabel('Staleness')

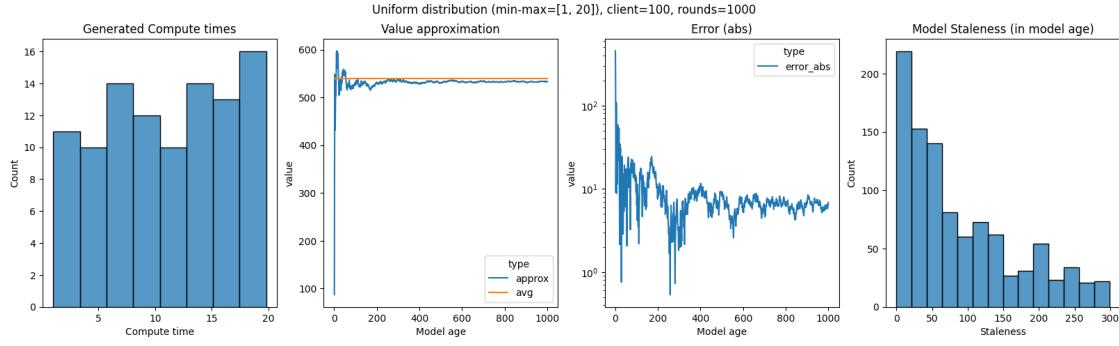
```

0.4 Generate results

```
[ ]: # Run with default values
# Num_clients = 1000
# Num_round = 10000
run()
```

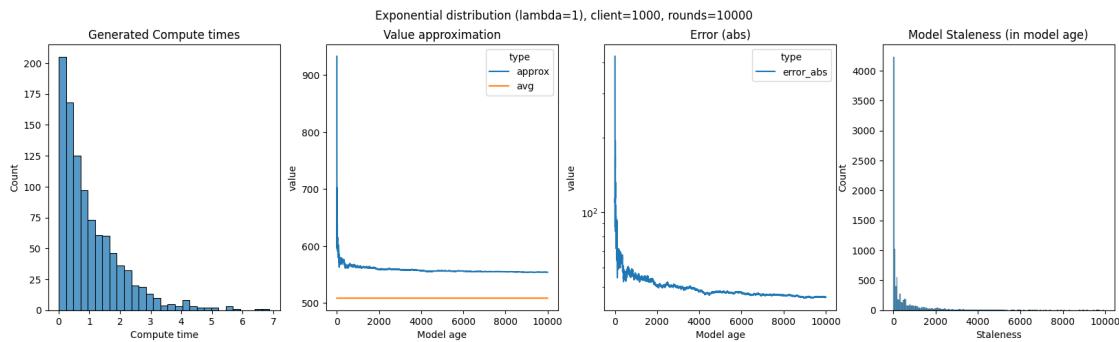
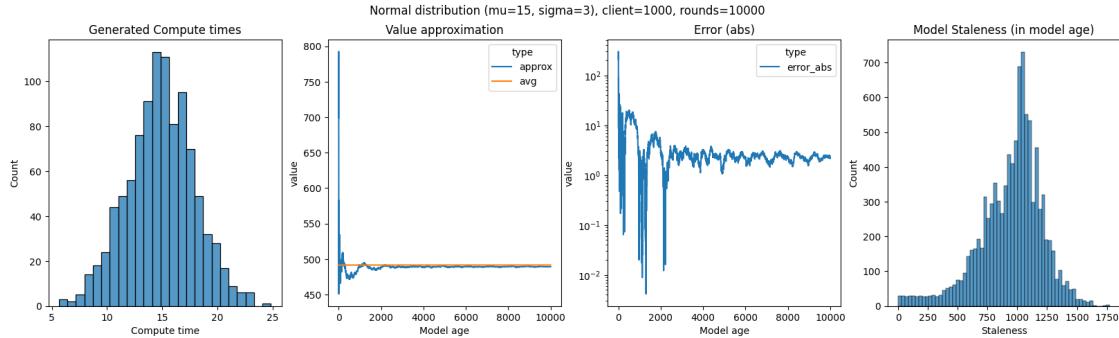
100% | 1000/1000 [00:00<00:00, 4174.27it/s]
100% | 1000/1000 [00:00<00:00, 3479.11it/s]
100% | 1000/1000 [00:00<00:00, 4198.22it/s]

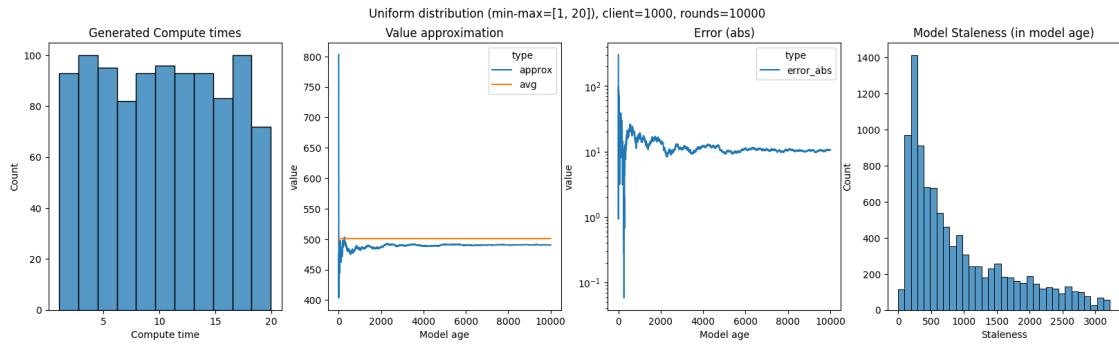




```
[ ]: # Run with large number of clients
run(num_clients=1000, num_rounds=10000)
```

100% | 10000/10000 [00:20<00:00, 498.84it/s]
 100% | 10000/10000 [00:20<00:00, 496.12it/s]
 100% | 10000/10000 [00:21<00:00, 459.56it/s]





[] :